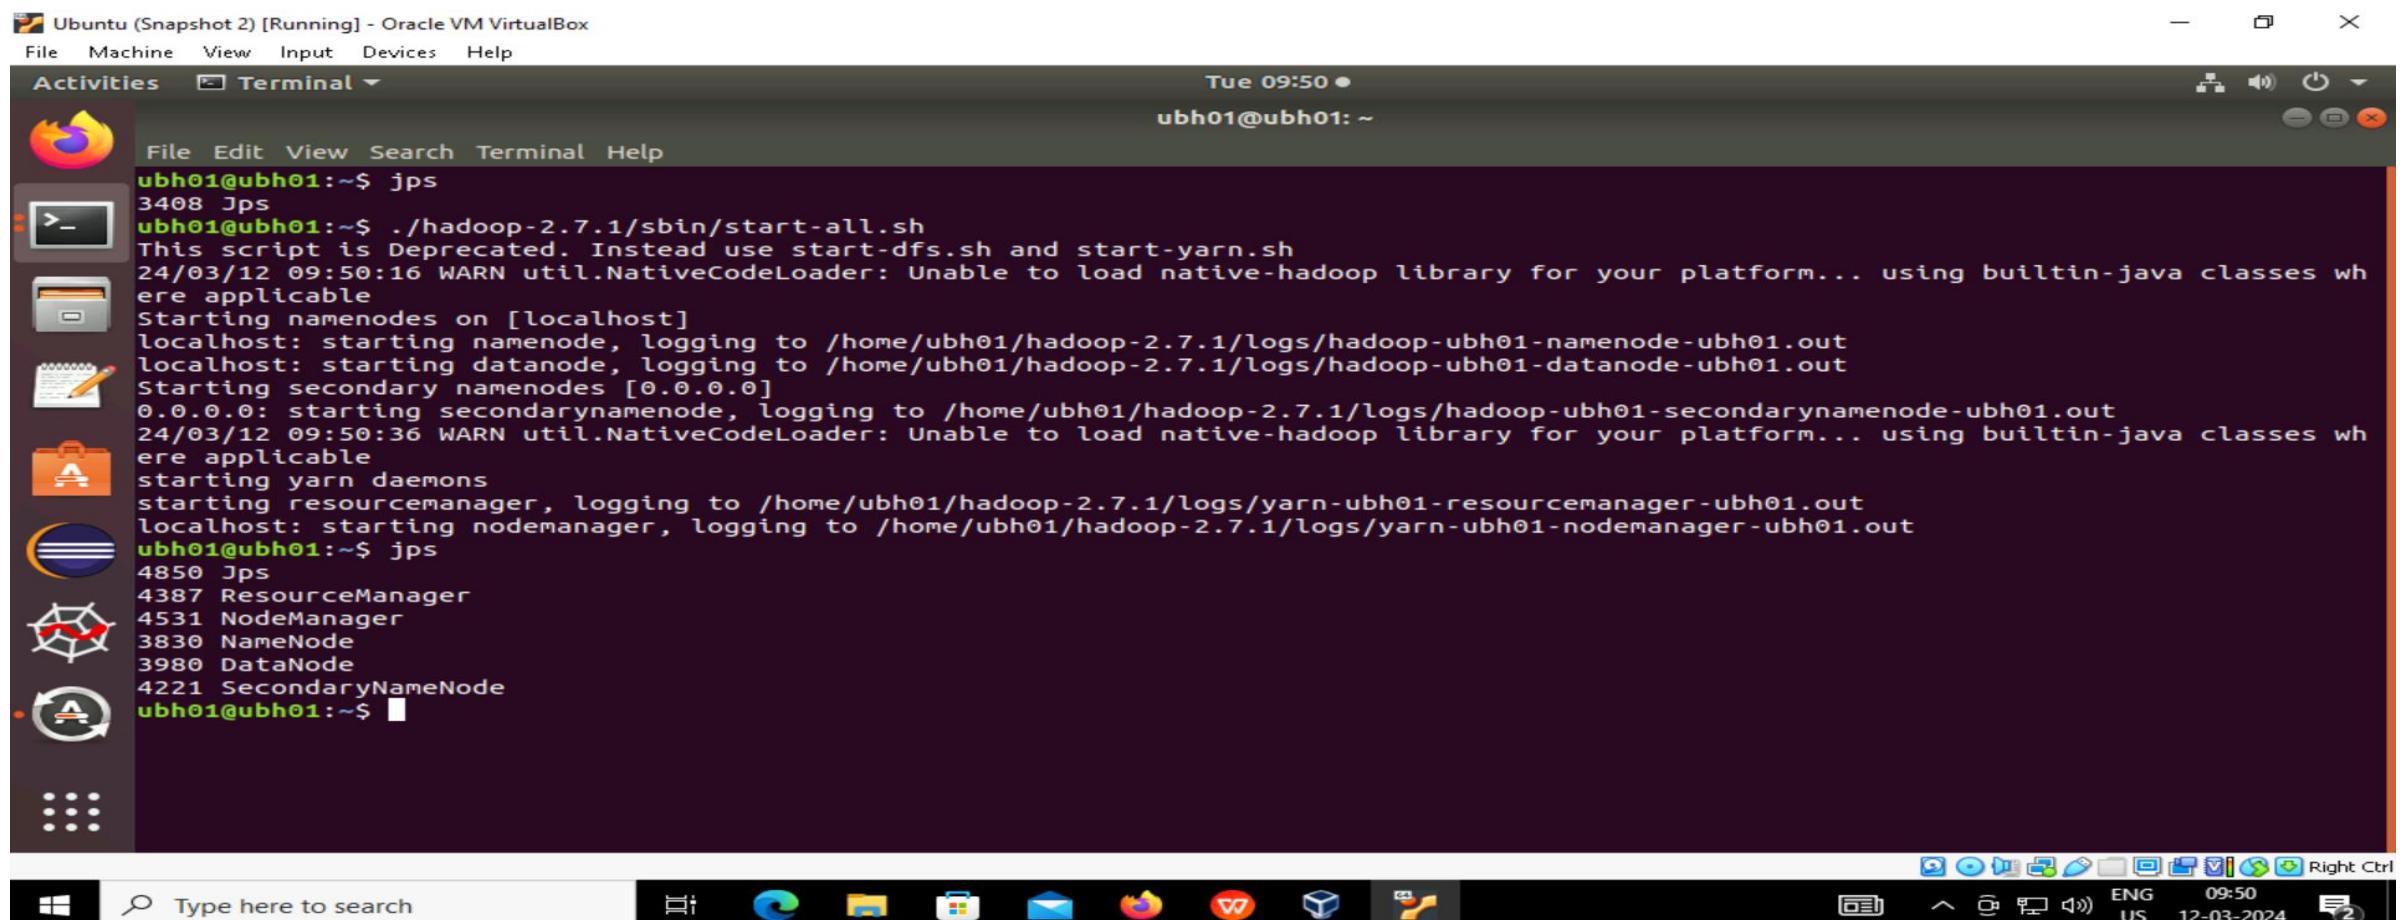


Employee Id : 2320393

Name : Sneha Ramachandran

RUN HADOOP

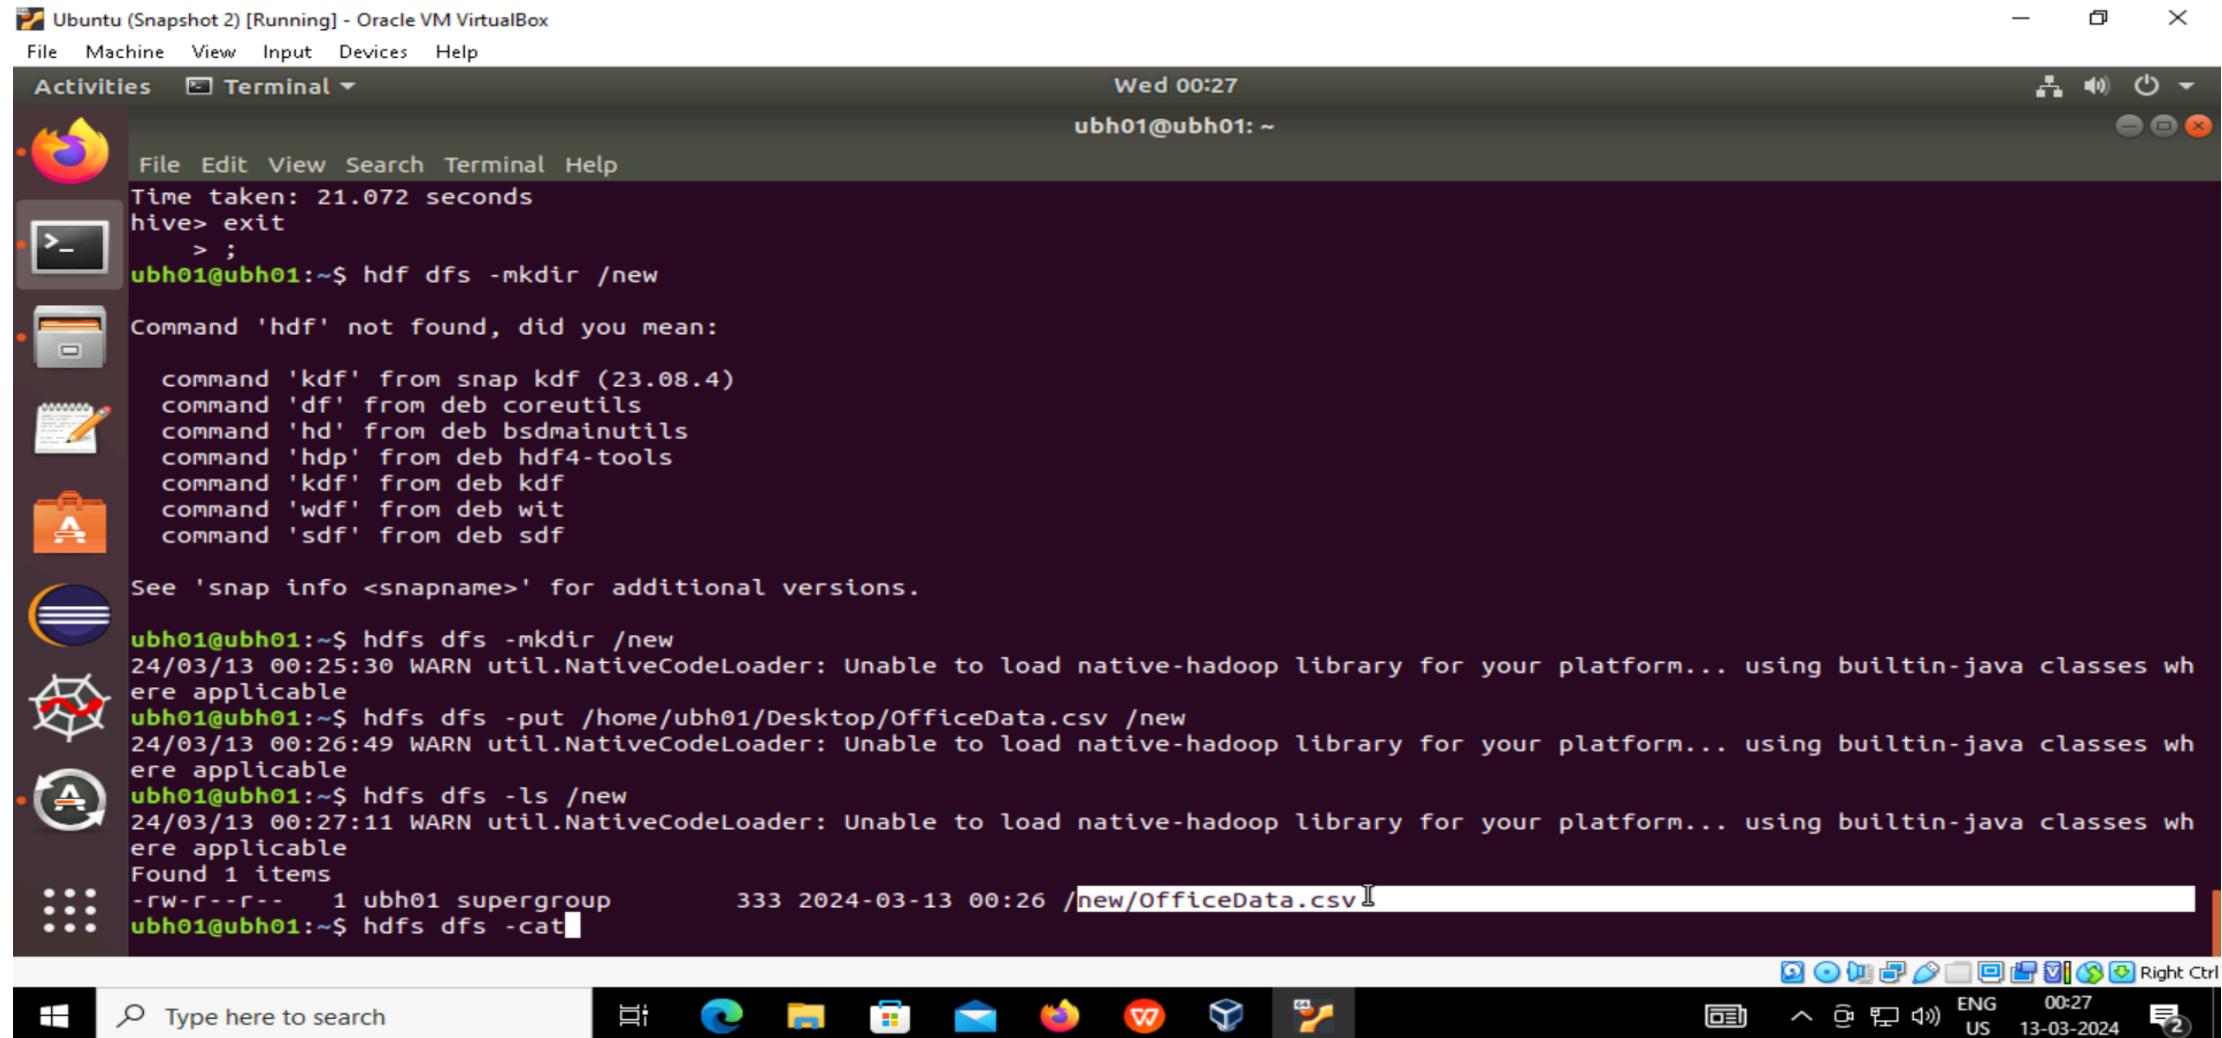


The screenshot shows a desktop environment with a terminal window open. The terminal window title is "Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox". The terminal content shows the execution of Hadoop startup scripts and a jps command to list Java processes.

```
ubh01@ubh01:~$ ./hadoop-2.7.1/sbin/start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
24/03/12 09:50:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/ubh01/hadoop-2.7.1/logs/hadoop-ubh01-namenode-ubh01.out
localhost: starting datanode, logging to /home/ubh01/hadoop-2.7.1/logs/hadoop-ubh01-datanode-ubh01.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/ubh01/hadoop-2.7.1/logs/hadoop-ubh01-secondarynamenode-ubh01.out
24/03/12 09:50:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /home/ubh01/hadoop-2.7.1/logs/yarn-ubh01-resourcemanager-ubh01.out
localhost: starting nodemanager, logging to /home/ubh01/hadoop-2.7.1/logs/yarn-ubh01-nodemanager-ubh01.out
ubh01@ubh01:~$ jps
4850 Jps
4387 ResourceManager
4531 NodeManager
3830 NameNode
3980 DataNode
4221 SecondaryNameNode
```

The desktop interface includes a dock at the bottom with icons for File Explorer, Mail, and a browser. The system tray shows network status, battery level, and system icons. The taskbar at the bottom right shows the date and time (12-03-2024, 09:50), language (ENG US), and a notification icon.

CREATE NEW DIRECTORY AND PUT THE LOCAL FILE TO HADOOP



Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal

Wed 00:27

ubh01@ubh01: ~

```
File Edit View Search Terminal Help
Time taken: 21.072 seconds
hive> exit
> ;
ubh01@ubh01:~$ hdf dfs -mkdir /new
Command 'hdf' not found, did you mean:
  command 'kdf' from snap kdf (23.08.4)
  command 'df' from deb coreutils
  command 'hd' from deb bsdmainutils
  command 'hdp' from deb hdf4-tools
  command 'kdf' from deb kdf
  command 'wdf' from deb wit
  command 'sdf' from deb sdf
See 'snap info <snapname>' for additional versions.

ubh01@ubh01:~$ hdfs dfs -mkdir /new
24/03/13 00:25:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
ubh01@ubh01:~$ hdfs dfs -put /home/ubh01/Desktop/OfficeData.csv /new
24/03/13 00:26:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
ubh01@ubh01:~$ hdfs dfs -ls /new
24/03/13 00:27:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 ubh01 supergroup          333 2024-03-13 00:26 /new/OfficeData.csv
ubh01@ubh01:~$ hdfs dfs -cat
```

Type here to search

Right Ctrl

ENG US 00:27 13-03-2024

DISPLAY THE FILE LOCATED IN HADOOP

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal

Wed 00:29

ubh01@ubh01: ~

```
ubh01@ubh01:~$ hdfs dfs -mkdir /new
24/03/13 00:25:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
ubh01@ubh01:~$ hdfs dfs -put /home/ubh01/Desktop/OfficeData.csv /new
24/03/13 00:26:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
ubh01@ubh01:~$ hdfs dfs -ls /new
24/03/13 00:27:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 ubh01 supergroup 333 2024-03-13 00:26 /new/OfficeData.csv
ubh01@ubh01:~$ hdfs dfs -cat /OfficeData.csv
24/03/13 00:28:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
cat: `/OfficeData.csv': No such file or directory
ubh01@ubh01:~$ hdfs dfs -cat /new/OfficeData.csv
24/03/13 00:28:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
employee_name,department,state,salary,age,bonus
James,Sales,NY,90000,34,10000
Michael,Sales,NY,86000,56,20000
Robert,Sales,CA,81000,30,23000
Maria,Finance,CA,90000,24,23000
Raman,Finance,CA,99000,40,24000
Scott,Finance,NY,83000,36,19000
Jen,Finance,NY,79000,53,15000
Jeff,Marketing,CA,80000,25,18000
Kumar,Marketing,NY,91000,50,21000
ubh01@ubh01:~$
```

Type here to search

Right Ctrl

ENG US 00:29 13-03-2024

RUN MYSQL

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Tue 09:59

ubh01@ubh01: ~

```
File Edit View Search Terminal Help
4221_SecondaryNameNode
ubh01@ubh01:~$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 2
Server version: 5.7.42-0ubuntu0.18.04.1 (Ubuntu)

Copyright (c) 2000, 2023, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> create database student;
Query OK, 1 row affected (0.02 sec)

mysql> use student;
Database changed
mysql> create table studentdata(sno int primary key,name varchar(100),age int,dept varchar(20));
Query OK, 0 rows affected (1.19 sec)

mysql> insert into studentdata(sno,name,age,dept) values (1,'sneha',21,'ECE'),(2,'thiru',21,'ECE'),(3,'dharni',22,'EEE'),(4,'sanjeev',23,'CSE');
Query OK, 4 rows affected (0.20 sec)
Records: 4 Duplicates: 0 Warnings: 0

mysql> select * from studendata;
ERROR 1146 (42S02): Table 'student.studendata' doesn't exist
```

Type here to search

Right Ctrl

ENG US 09:59 12-03-2024

DISPLAY TABLE AND EXIT FROM MYSQL

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Tue 10:00

ubh01@ubh01: ~

```
File Edit View Search Terminal Help
mysql> create database student;
Query OK, 1 row affected (0.02 sec)

mysql> use student;
Database changed
mysql> create table studentdata(sno int primary key,name varchar(100),age int,dept varchar(20));
Query OK, 0 rows affected (1.19 sec)

mysql> insert into studentdata(sno,name,age,dept) values (1,'sneha',21,'ECE'),(2,'thiru',21,'ECE'),(3,'dharni',22,'EEE'),(4,'sanjeev',23,'CSE');
Query OK, 4 rows affected (0.20 sec)
Records: 4  Duplicates: 0  Warnings: 0

mysql> select * from studendata;
ERROR 1146 (42S02): Table 'student.studendata' doesn't exist
mysql> select * from studentdata;
+----+-----+-----+
| sno | name  | age   | dept  |
+----+-----+-----+
| 1  | sneha | 21    | ECE   |
| 2  | thiru | 21    | ECE   |
| 3  | dharni| 22    | EEE   |
| 4  | sanjeev| 23    | CSE   |
+----+-----+-----+
4 rows in set (0.01 sec)

mysql> exit
Bye
ubh01@ubh01:~$
```

Type here to search

Right Ctrl

ENG US 10:00 12-03-2024

IMPORT SQOOP

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Tue 10:06

ubh01@ubh01: ~

```
mysql> exit
Bye
ubh01@ubh01:~$ sqoop import --connect jdbc:mysql://localhost:3306/student --username root --password password --table studentdata --target-dir /s1
Warning: /home/ubh01/sqoop-1.4.7-bin_hadoop2.6.0//../.hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/ubh01/sqoop-1.4.7-bin_hadoop2.6.0//../.accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /home/ubh01/sqoop-1.4.7-bin_hadoop2.6.0//../.zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
24/03/12 10:04:58 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
24/03/12 10:04:58 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
24/03/12 10:04:58 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
24/03/12 10:04:58 INFO tool.CodeGenTool: Beginning code generation
Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
24/03/12 10:04:59 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `studentdata` AS t LIMIT 1
24/03/12 10:04:59 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `studentdata` AS t LIMIT 1
24/03/12 10:04:59 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/ubh01/hadoop-2.7.1
Note: /tmp/sqoop-ubh01/compile/ef72bfce73af93d4fb9b1e7f5520f639/studentdata.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
24/03/12 10:05:04 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-ubh01/compile/ef72bfce73af93d4fb9b1e7f5520f639/studentdata.jar
24/03/12 10:05:04 WARN manager.MySQLManager: It looks like you are importing from mysql.
24/03/12 10:05:04 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
24/03/12 10:05:04 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
24/03/12 10:05:04 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
24/03/12 10:05:04 INFO mapreduce.ImportJobBase: Beginning import of studentdata
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in file:/home/ubh01/hadoop-2.7.1/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLogger
```

Type here to search

Right Ctrl

ENG 10:06
US 12-03-2024

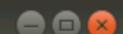
File Machine View Input Devices Help

Activities Terminal

Tue 10:06



ubh01@ubh01: ~



File Edit View Search Terminal Help

```
HDFS: Number of large read operations=0  
HDFS: Number of write operations=8
```

Job Counters

```
Launched map tasks=4  
Other local map tasks=4  
Total time spent by all maps in occupied slots (ms)=122807  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=122807  
Total vcore-seconds taken by all map tasks=122807  
Total megabyte-seconds taken by all map tasks=125754368
```

Map-Reduce Framework

```
Map input records=4  
Map output records=4  
Input split bytes=401  
Spilled Records=0  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=1958  
CPU time spent (ms)=13950  
Physical memory (bytes) snapshot=713850880  
Virtual memory (bytes) snapshot=7694290944  
Total committed heap usage (bytes)=548405248
```

File Input Format Counters

```
Bytes Read=0
```

File Output Format Counters

```
Bytes Written=63
```

```
24/03/12 10:06:11 INFO mapreduce.ImportJobBase: Transferred 63 bytes in 65.3976 seconds (0.9633 bytes/sec)
```

```
24/03/12 10:06:11 INFO mapreduce.ImportJobBase: Retrieved 4 records.
```

```
ubh01@ubh01:~$ █
```



Type here to search

Right Ctrl
ENG US 10:06 12-03-2024

O/P OF SQOOP IMPORT

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Tue 10:08 ubh01@ubh01:~

```
Merged Map outputs=0
GC time elapsed (ms)=1958
CPU time spent (ms)=13950
Physical memory (bytes) snapshot=713850880
Virtual memory (bytes) snapshot=7694290944
Total committed heap usage (bytes)=548405248
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=63
24/03/12 10:06:11 INFO mapreduce.ImportJobBase: Transferred 63 bytes in 65.3976 seconds (0.9633 bytes/sec)
24/03/12 10:06:11 INFO mapreduce.ImportJobBase: Retrieved 4 records.
ubh01@ubh01:~$ hdfs dfs -ls /s1
24/03/12 10:07:04 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 5 items
-rw-r--r-- 1 ubh01 supergroup 0 2024-03-12 10:06 /s1/_SUCCESS
-rw-r--r-- 1 ubh01 supergroup 15 2024-03-12 10:05 /s1/part-m-00000
-rw-r--r-- 1 ubh01 supergroup 15 2024-03-12 10:05 /s1/part-m-00001
-rw-r--r-- 1 ubh01 supergroup 16 2024-03-12 10:05 /s1/part-m-00002
-rw-r--r-- 1 ubh01 supergroup 17 2024-03-12 10:05 /s1/part-m-00003
ubh01@ubh01:~$ hdfs dfs -cat /s1/part*
24/03/12 10:08:15 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
1,sneha,21,ECE
2,thiru,21,ECE
3,dharni,22,EEE
4,sanjeev,23,CSE
ubh01@ubh01:~$
```

Type here to search

Right Ctrl

ENG US 10:08 12-03-2024

EXPORT SQOOP

```
v Input Devices Help
ubh01@ubh01:~$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 17
Server version: 5.7.36-0ubuntu0.18.04.1 (Ubuntu)

Copyright (c) 2000, 2021, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

?
mysqldb> use samp;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
I
Database changed
mysqldb> create table c1(name varchar(100),id int,age int,disease varchar(10));
Query OK, 0 rows affected (0.01 sec)

mysqldb> exit
Bye
ubh01@ubh01:~$ sqoop export --connect jdbc:mysql://localhost:3306/samp --username root --password password --table c1 --export-dir /s
Warning: /home/ubh01/sqoop-1.4.7-bin_hadoop-2.6.0/..../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/ubh01/sqoop-1.4.7-bin_hadoop-2.6.0/..../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /home/ubh01/sqoop-1.4.7-bin_hadoop-2.6.0/..../zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
24/03/11 20:58:16 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
24/03/11 20:58:16 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
24/03/11 20:58:16 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
24/03/11 20:58:16 INFO tool.CodeGenTool: Beginning code generation
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
24/03/11 20:58:17 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `c1` AS t LIMIT 1
24/03/11 20:58:17 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `c1` AS t LIMIT 1
24/03/11 20:58:17 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/ubh01/hadoop-2.7.1
Note: /tmp/sqoop-ubh01/compile/f070b4d0555f1bf7f1a9bd9b02a97bc4/c1.java uses or overrides a deprecated API.
```

```
CPUTIME spent (ms)=5160
Physical memory (bytes) snapshot=711806976
Virtual memory (bytes) snapshot=7751335936
Total committed heap usage (bytes)=561512448
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=0
24/03/11 20:58:33 INFO mapreduce.ExportJobBase: Transferred 713 bytes in 14.9079 seconds (47.8271 bytes/sec)
24/03/11 20:58:33 INFO mapreduce.ExportJobBase: Exported 4 records.
ubh01@ubh01:~$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 23
Server version: 5.7.36-0ubuntu0.18.04.1 (Ubuntu)

Copyright (c) 2000, 2021, Oracle and/or its affiliates.           I

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> use samp;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> select * from c1;
+-----+-----+-----+
| name | id  | age  | disease |
+-----+-----+-----+
| san  | 46  | 30   | cold    |
| sne  | 39  | 21   | ulcer   |
| veni | 36  | 22   | fever   |
| dharni | 29 | 23   | appe    |
+-----+-----+-----+
4 rows in set (0.00 sec)

mysql>
```

O/P OF SQOOP EXPORT

The screenshot shows a terminal window on an Ubuntu desktop. The terminal output is as follows:

```
Tue 10:21  
ubh01@ubh01:~  
File Edit View Terminal Help  
24/03/12 10:15:41 INFO mapreduce.ImportJobBase: Retrieved 2 records.  
ubh01@ubh01:~$ dhfs dfs -cat /s2/part*  
Command 'dhfs' not found, did you mean:  
  command 'hdfs' from deb hfsutils-tcltk  
  command 'xhdfs' from deb hfsutils-tcltk  
Try: sudo apt install <deb name>  
ubh01@ubh01:~$ hdfs dfs -cat /s2/part*  
24/03/12 10:16:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
1,sneha,21,ECE  
2,thiru,21,ECE  
ubh01@ubh01:~$ hdfs dfs -cat /s1/part*  
24/03/12 10:20:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
1,sneha,21,ECE  
2,thiru,21,ECE  
3,dharni,22,EEE  
4,sanjeev,23,CSE  
ubh01@ubh01:~$ sqoop import --connect jdbc:mysql://localhost:3306/student --username root --password password --table studentdata --target-dir /s1  
Warning: /home/ubh01/sqoop-1.4.7.bin__hadoop-2.6.0/..../hcatalog does not exist! HCatalog jobs will fail.  
Please set $HCAT_HOME to the root of your HCatalog installation.  
Warning: /home/ubh01/sqoop-1.4.7.bin__hadoop-2.6.0/..../accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
Warning: /home/ubh01/sqoop-1.4.7.bin__hadoop-2.6.0/..../zookeeper does not exist! Accumulo imports will fail.  
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
```

SPLIT-BY IN SQOOP

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Tue 10:21

ubh01@ubh01: ~

```
ubh01@ubh01:~$ sqoop import --connect jdbc:mysql://localhost:3306/student --username root --password password --query "select * from studentdata where dept='ECE' and \$CONDITIONS" --split-by sno --target-dir /s2
Warning: /home/ubh01/sqoop-1.4.7.bin__hadoop-2.6.0//..hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/ubh01/sqoop-1.4.7.bin__hadoop-2.6.0//..accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /home/ubh01/sqoop-1.4.7.bin__hadoop-2.6.0//..zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
24/03/12 10:15:05 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
24/03/12 10:15:05 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
24/03/12 10:15:05 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
24/03/12 10:15:05 INFO tool.CodeGenTool: Beginning code generation
Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
24/03/12 10:15:07 INFO manager.SqlManager: Executing SQL statement: select * from studentdata where dept='ECE' and (1 = 0)
24/03/12 10:15:07 INFO manager.SqlManager: Executing SQL statement: select * from studentdata where dept='ECE' and (1 = 0)
24/03/12 10:15:07 INFO manager.SqlManager: Executing SQL statement: select * from studentdata where dept='ECE' and (1 = 0)
24/03/12 10:15:07 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/ubh01/hadoop-2.7.1
Note: /tmp/sqoop-ubh01/compile/8bcba78e00d52e3c25cef59c59a0c3c0/QueryResult.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
24/03/12 10:15:10 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-ubh01/compile/8bcba78e00d52e3c25cef59c59a0c3c0/QueryResult.jar
24/03/12 10:15:10 INFO mapreduce.ImportJobBase: Beginning query import.
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/ubh01/hadoop-2.7.1/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/ubh01/hbase-1.1.2/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
24/03/12 10:15:10 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes wh...
```

Type here to search

Right Ctrl

ENG 10:21
US 12-03-2024

RUN HIVE

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Tue 23:55
ubh01@ubh01: ~

```
File Edit View Search Terminal Help
5420 SecondaryNameNode
ubh01@ubh01:~$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/ubh01/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/ubh01/hadoop-2.7.1/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/ubh01/apache-hive-2.3.2-bin/lib/hive-common-2.3.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> create table office(employee_name varchar(100),department varchar(100),state varchar(100),salary int,age int,bonus int) row format delimited fields terminated by ','
      > tblproperties ('skip.header.line.count='1');
Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
OK
Time taken: 6.806 seconds
hive> load data local inpath '/home/ubh01/Desktop/OfficeData.csv' into table office;
Loading data to table default.office
OK
Time taken: 1.98 seconds
hive> load data local inpath '/home/ubh01/Desktop/hive/OfficeData.csv' into table office;
FAILED: SemanticException Line 1:23 Invalid path ''/home/ubh01/Desktop/hive/OfficeData.csv'': No files matching path file:/home/ubh01/Desktop/hive/OfficeData.csv
hive>
```

Type here to search

Right Ctrl

File Machine View Input Devices Help

Activities Terminal Tue 23:55
ubh01@ubh01: ~

File Edit View Search Terminal Help

5420 SecondaryNameNode

ubh01@ubh01:~\$ hive

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/ubh01/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/ubh01/hadoop-2.7.1/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/ubh01/apache-hive-2.3.2-bin/lib/hive-common-2.3.2.jar!/hive-log4j2.properties

Async: true

Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

hive> create table office(employee_name varchar(100),department varchar(100),state varchar(100),salary int,age int,bonus int) row format delimited fields terminated by ','

>tblproperties ('skip.header.line.count='1');

Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.

OK

Time taken: 6.806 seconds

hive> load data local inpath '/home/ubh01/Desktop/OfficeData.csv' into table office;

Loading data to table default.office

OK

Time taken: 1.98 seconds

hive> load data local inpath '/home/ubh01/Desktop/hive/OfficeData.csv' into table office;

FAILED: SemanticException Line 1:23 Invalid path ''/home/ubh01/Desktop/hive/OfficeData.csv'': No files matching path file:/home/ubh01/Desktop/hive/OfficeData.csv

hive>

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal

Tue 23:57

ubh01@ubh01: ~



File Edit View Search Terminal Help

```
at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1457)
at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1237)
at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1227)
at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:233)
at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:184)
at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:403)
at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:821)
at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:759)
at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:686)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
```

```
FAILED: ParseException line 1:20 cannot recognize input near 'office' ':' 'select' in joinSource
hive> select * from office;
```

OK

James	Sales	NY	90000	34	10000
Michael	Sales	NY	86000	56	20000
Robert	Sales	CA	81000	30	23000
Maria	Finance	CA	90000	24	23000
Raman	Finance	CA	99000	40	24000
Scott	Finance	NY	83000	36	19000
Jen	Finance	NY	79000	53	15000
Jeff	Marketing	CA	80000	25	18000
Kumar	Marketing	NY	91000	50	21000

Time taken: 2.106 seconds, Fetched: 9 row(s)

hive> create



Right Ctrl
Type here to search ENG 23:57
US 12-03-2024

STATIC PARTITIONING

INSERT INTO STATIC PARTITION TABLE

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Wed 00:07
ubh01@ubh01: ~

File Edit View Search Terminal Help

```
at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:686)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
FAILED: ParseException line 1:130 cannot recognize input near 'Sales' ';' '<EOF>' in expression specification
hive> insert into table sp partition(department='Sales') select employee_name,state,salary,age,bonus from office where department='Sales';
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = ubh01_20240313000522_d23b8c9b-ed33-4ac8-b149-28db61f89c80
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1710267179929_0001, Tracking URL = http://ubh01:8088/proxy/application_1710267179929_0001/
Kill Command = /home/ubh01/hadoop-2.7.1/bin/hadoop job -kill job_1710267179929_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-03-13 00:05:48,219 Stage-1 map = 0%, reduce = 0%
2024-03-13 00:05:54,725 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.7 sec
MapReduce Total cumulative CPU time: 3 seconds 700 msec
Ended Job = job_1710267179929_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/sp/department=Sales/.hive-staging_hive_2024-03-13_00-05-22_036_333
7032557709211843-1/-ext-10000
Loading data to table default.sp partition (department=Sales)
MapReduce Jobs Launched:
```

Type here to search ENG US 00:07 13-03-2024

O/P OF STATIC PARTITIONING

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal ▾

Wed 00:07

ubh01@ubh01: ~

```
Query ID = ubh01_20240313000522_d23b8c9b-ed33-4ac8-b149-28db61f89c80
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1710267179929_0001, Tracking URL = http://ubh01:8088/proxy/application_1710267179929_0001/
Kill Command = /home/ubh01/hadoop-2.7.1/bin/hadoop job -kill job_1710267179929_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-03-13 00:05:48,219 Stage-1 map = 0%, reduce = 0%
2024-03-13 00:05:54,725 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.7 sec
MapReduce Total cumulative CPU time: 3 seconds 700 msec
Ended Job = job_1710267179929_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/sp/department=Sales/.hive-staging_hive_2024-03-13_00-05-22_036_333
7032557709211843-1/-ext-10000
Loading data to table default.sp partition (department=Sales)
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 3.7 sec   HDFS Read: 5450 HDFS Write: 158 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 700 msec
OK
Time taken: 35.979 seconds
hive> select * from sp;
OK
James    NY      90000    34      10000    Sales
Michael  NY      86000    56      20000    Sales
Robert   CA      81000    30      23000    Sales
Time taken: 0.312 seconds, Fetched: 3 row(s)
hive>
```

Type here to search

Right Ctrl

File Home App Drawer Dash Search Mail Browser Terminal File Home Dash Search Mail Browser Terminal

ENG US 00:07 13-03-2024

DYNAMIC PARTITIONING

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Wed 00:12
ubh01@ubh01: ~

```
hive> create table dp(employee_name varchar(100),state varchar(100),salary int,age int,bonus int) partitioned by(department varchar(100));
OK
Time taken: 1.106 seconds
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> insert into table sp partition(department) select employee_name,state,salary,age,bonus,department from office;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = ubh01_20240313001145_402968f1-a810-4d9e-a992-cd11011425a7
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1710267179929_0002, Tracking URL = http://ubh01:8088/proxy/application_1710267179929_0002/
Kill Command = /home/ubh01/hadoop-2.7.1/bin/hadoop job -kill job_1710267179929_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-03-13 00:11:53,038 Stage-1 map = 0%, reduce = 0%
2024-03-13 00:12:04,719 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.58 sec
MapReduce Total cumulative CPU time: 3 seconds 580 msec
Ended Job = job_1710267179929_0002
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/sp/.hive-staging_hive_2024-03-13_00-11-45_338_2297604718206082014-1/-ext-10000
Loading data to table default.sp partition (department=null)

Loaded : 3/3 partitions.
Time taken to load dynamic partitions: 1.204 seconds
Time taken for adding to write entity : 0.0 seconds
MapReduce Jobs Launched.
```

Type here to search Right Ctrl

File Machine View Input Devices Help

Activities Terminal Wed 00:12
ubh01@ubh01: ~

File Edit View Search Terminal Help

hive> create table dp(employee_name varchar(100),state varchar(100),salary int,age int,bonus int) partitioned by(department varchar(100));

OK

Time taken: 1.106 seconds

hive> set hive.exec.dynamic.partition.mode=nonstrict;

hive> insert into table sp partition(department) select employee_name,state,salary,age,bonus,department from office;

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = ubh01_20240313001145_402968f1-a810-4d9e-a992-cd11011425a7

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1710267179929_0002, Tracking URL = http://ubh01:8088/proxy/application_1710267179929_0002/

Kill Command = /home/ubh01/hadoop-2.7.1/bin/hadoop job -kill job_1710267179929_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2024-03-13 00:11:53,038 Stage-1 map = 0%, reduce = 0%

2024-03-13 00:12:04,719 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.58 sec

MapReduce Total cumulative CPU time: 3 seconds 580 msec

Ended Job = job_1710267179929_0002

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/sp/.hive-staging_hive_2024-03-13_00-11-45_338_2297604718206082014-1/-ext-10000

Loading data to table default.sp partition (department=null)

Loaded : 3/3 partitions.

Time taken to load dynamic partitions: 1.204 seconds

Time taken for adding to write entity : 0.0 seconds

MapReduce Jobs Launched.



File Edit View Search Terminal Help

```
Time taken: 0.518 seconds
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> insert into table dp partition(Gender) select Year,FirstName,LastName,Age,Gender from baby3;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = ubh01_20240312224836_cbc57b1-01a9-4229-8566-cfba554c173d
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1710261214077_0003, Tracking URL = http://ubh01:8088/proxy/application_1710261214077_0003/
Kill Command = /home/ubh01/hadoop-2.7.1/bin/hadoop job -kill job_1710261214077_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-03-12 22:48:42,241 Stage-1 map = 0%, reduce = 0%
2024-03-12 22:48:47,489 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.49 sec
MapReduce Total cumulative CPU time: 2 seconds 490 msec
Ended Job = job_1710261214077_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/dp/.hive-staging_hive_2024-03-12_22-48-36_570_1133809601253564079-1/-ext-10000
Loading data to table default.dp partition (gender=null)

Loaded : 2/2 partitions.
    Time taken to load dynamic partitions: 0.402 seconds
    Time taken for adding to write entity : 0.0 seconds
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 2.49 sec   HDFS Read: 1321186 HDFS Write: 1159861 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 490 msec
OK
```

Right Ctrl



Type here to search

ENG
US 22:49
12-03-2024

BUCKETING

Ubuntu (Snapshot 2) [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Thu 10:02 ubh01@ubh01: ~

```
File Edit View Search Terminal Help
2014    JACKSON Wyoming M      5
Time taken: 2.477 seconds, Fetched: 52252 row(s)
hive> create external table bby1(Year int,FirstName varchar(100),LastName varchar(100),Gender varchar(100),Age int)
      > row format delimited fields terminated by ','
      > tblproperties('skip.header.line.count='1');
OK
Time taken: 1.021 seconds
hive> load data local inpath '/home/ubh01/Desktop/baby_names.csv' into table bby1;
Loading data to table default.bby1
OK
Time taken: 0.803 seconds
hive> select * from bby1 limit 5;
OK
2013    GAVIN    ST LAWRENCE    M      9
2013    LEVI     ST LAWRENCE    M      9
2013    LOGAN    NEW YORK      M     44
2013    HUDSON   NEW YORK      M     49
2013    GABRIEL  NEW YORK      M     50
Time taken: 0.459 seconds, Fetched: 5 row(s)
hive> create table bk(Year int,FirstName varchar(100),LastName varchar(100),Gender varchar(100),Age int) clustered by(Year) into 3 buckets;
OK
Time taken: 0.881 seconds
hive> insert into bk select * from bby1;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = ubh01_20240314095519_19014b94-eda6-4823-9c37-503bb173d997
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
```

Type here to search    Right Ctrl

File Machine View Input Devices Help

Activities Terminal Thu 10:02 ubh01@ubh01: ~

File Edit View Search Terminal Help

2014 JACKSON Wyoming M 5

Time taken: 2.477 seconds, Fetched: 52252 row(s)

hive> create external table bby1(Year int,FirstName varchar(100),LastName varchar(100),Gender varchar(100),Age int)

> row format delimited fields terminated by ','

>tblproperties('skip.header.line.count='1');

OK

Time taken: 1.021 seconds

hive> load data local inpath '/home/ubh01/Desktop/baby_names.csv' into table bby1;

Loading data to table default.bby1

OK

Time taken: 0.803 seconds

hive> select * from bby1 limit 5;

OK

Year	FirstName	LastName	Gender	Age
2013	GAVIN	ST LAWRENCE	M	9
2013	LEVI	ST LAWRENCE	M	9
2013	LOGAN	NEW YORK	M	44
2013	HUDSON	NEW YORK	M	49
2013	GABRIEL	NEW YORK	M	50

Time taken: 0.459 seconds, Fetched: 5 row(s)

hive> create table bk(Year int,FirstName varchar(100),LastName varchar(100),Gender varchar(100),Age int) clustered by(Year) into 3 buckets;

OK

Time taken: 0.881 seconds

hive> insert into bk select * from bby1;

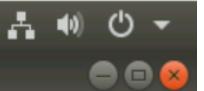
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = ubh01_20240314095519_19014b94-eda6-4823-9c37-503bb173d997

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 3



File Edit View Search Terminal Help

2007 INDAUFFOLK
2007 INDAQUEENS
2007 INDSAYUFFOLK
2007 INDSAYASSAU
2007 INDSAYDNROE

ubh01@ubh01:~\$ hdfs dfs -ls /

24/03/14 10:06:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
^[[A^[[AFound 2 itemsdrwx-wx-wx - ubh01 supergroup 0 2024-03-14 09:55 /tmp
drwxr-xr-x - ubh01 supergroup 0 2022-01-20 17:54 /user

ubh01@ubh01:~\$ hdfs dfs -ls /user/hive/warehouse/bk

24/03/14 10:06:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items-rwxrwxr-x 1 ubh01 supergroup 454348 2024-03-14 09:55 /user/hive/warehouse/bk/000000_0
-rwxrwxr-x 1 ubh01 supergroup 507449 2024-03-14 09:55 /user/hive/warehouse/bk/000001_0
-rwxrwxr-x 1 ubh01 supergroup 302441 2024-03-14 09:55 /user/hive/warehouse/bk/000002_0

ubh01@ubh01:~\$ hdfs dfs -ls /user/hive/warehouse/bby1

24/03/14 10:07:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items

-rwxrwxr-x 1 ubh01 supergroup 1316524 2024-03-14 09:52 /user/hive/warehouse/bby1/baby_names.csv

ubh01@ubh01:~\$ hdfs dfs -ls /tmp

24/03/14 10:07:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 itemsdrwx----- - ubh01 supergroup 0 2024-03-14 09:55 /tmp/hadoop-yarn
drwx-wx-wx - ubh01 supergroup 0 2022-01-20 18:25 /tmp/hive

ubh01@ubh01:~\$ hdfs dfs -ls /tmp/hadoop-yarn

24/03/14 10:08:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items

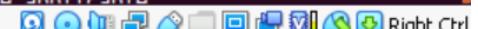
drwx----- - ubh01 supergroup 0 2024-03-14 09:55 /tmp/hadoop-yarn/staging

ubh01@ubh01:~\$ hdfs dfs -ls /tmp/hadoop-yarn/staging

24/03/14 10:08:27 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 itemsdrwxr-xr-x - ubh01 supergroup 0 2024-03-14 09:55 /tmp/hadoop-yarn/staging/history
drwx----- - ubh01 supergroup 0 2024-03-14 09:55 /tmp/hadoop-yarn/staging/ubh01

ubh01@ubh01:~\$ hdfs dfs -ls /tmp/hadoop-yarn/staging/history

24/03/14 10:08:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable



Type here to search

ENG
US 10:09
14-03-2024

2